

FDE Take-Home Review

Meridian Systems / 45 min / 5-10 min walkthrough + Q&A; / Hire · Lean Hire · No Hire

WHAT YOU'RE EVALUATING

- 1 Did they start with the business problem or their tech stack?
- 2 Did they build something real, or just a deck?
- 3 How do they handle the hallucination problem, specifically?

The job is embedding with clients who have messy, ambiguous problems and low tolerance for AI that makes things up. That is the filter.

THE WALKTHROUGH (5-10 MIN -- DO NOT INTERRUPT)

Listen for whether they open with the business problem (CCO trust, slow pipeline answers) or whether they open with their architecture.

They should demo the artifact live, not just describe it.

When they hit a limitation: do they own it or gloss past it?

Over 10 minutes: "Let's pause and move to questions."

ARTIFACT QUICK-SCAN (CHECK OFF AS THEY PRESENT)

STANDARD

- ☐ "How is Enterprise tracking against quota this quarter?"
- ☐ "Which reps are at risk of missing Q2?"
- ☐ "How does Q2 compare to Q1 at the same point?"

TRICKY

- ☐ REP-99 -- does not exist, should flag not invent
- ☐ EMEA pipeline -- not in dataset, should flag not invent
- ☐ "This quarter?" with no context -- should surface Q1/Q2 ambiguity

If none of the tricky ones work: "Walk me through what happens when I ask about REP-99."

PROBE QUESTIONS (PICK 2 OR 3)

On scope: "What's the first thing you'd cut if the interview was tomorrow morning?"

GOOD

Fast, specific. They have thought about this.

WEAK

"I would not cut anything" or a vague non-answer.

On ambiguous data: "Dataset has Q1 and Q2. User asks 'how are we doing this quarter?' -- what does your system do?"

GOOD

Surfaces the ambiguity. Q1 is closed and final. Q2 is 34 days in with no quota yet. No clean answer without knowing what the user needs.

WEAK

Picks one quarter and answers. Does not flag the ambiguity.

On production: "If 50 Meridian users started using this next week, what breaks first?"

GOOD

Specific: latency, cost, stale data, no auth, hallucination rate.

WEAK

"I would add more features first" or no concrete failure mode.

On Azure: "Where would Azure AI Foundry fit in if you had another week?"

GOOD

Knows it. Can name model gateway, prompt management, or eval framework.

WEAK

Blank look. Does not know what Azure AI Foundry is.

THE TRUST QUESTION

Raise this in Q&A; -- it may come up in the walkthrough or you may need to ask it directly:

"The last AI tool this client used lost executive trust because it hallucinated. How does your approach avoid that problem, and where does it still fall short?"

GOOD

Names specific mechanisms: citations, code execution to verify numbers, flags when data is missing.

Acknowledges real failure modes -- model can still be wrong, edge cases exist, data goes stale.

Talks about the skeptical CCO, not AI safety in the abstract.

WEAK

"Our model is accurate" or "I tested it."

No traceability mechanism named.

No honest accounting of where it still fails.